

Data manipulation in R

Surya Lamichhane

10 February 2025

Contents

Learning Objectives	1
Data Wrangling	1
1. Data Exploration in R	2
2. Subsetting Dataframes	5
3. Sorting data in R	9
4. Merging of Data Tables	11
5. dplyr Package for Data Manipulation	14
6. Data Reshaping using tidyr	19

Learning Objectives

By the end of this lesson, you will be able to:

- Understand basic data wrangling techniques
- Read and manipulate data using functions in R
- Summarize data
- Use dplyr for data manipulation
- Merge data tables to prepare data for analysis

Data Wrangling

- A data wrangling is a process of transforming and mapping data from “raw” form into more usable format
- It is used for a variety of proposes including analytics and decision making.

Data Wrangling Steps

- Discovery: It contains process of exploring the data and finding the way to solve the data analysis objective
- Structuring: The process of taking raw data and converting it into a more usable structured data that fits to the relevant study

- **Cleaning:** The process of removing errors that must be cleaned before it can be used. Data cleaning includes correcting outliers, deleting unreliable data to increase quality and consistency. It finds duplicate values, eliminates structural problems, and verifies data to make it easier to use.
- **Enriching:** This is a data augmentation process adding more information to the data from other datasets that might improve the present analysis.
- **Validating:** This is a process of ensuring that the processed data is accurate and consistent. Through this step you will ensure that your data is ready for analysis.
- **Publishing:** This is a step where you finalize the data and make it available to other stakeholders.

1. Data Exploration in R

Data exploration is a critical step in understanding the structure, quality, and patterns in your dataset before performing any analysis. In this course, we are performing basic and intermediate data exploration in R.

1.1. Reading data in R

A. Working directory: The working directory is a default location or path of any files to be read into R, or saved out of R.

- Get the current working directory: `getwd()`
- Set the new working directory: `setwd("/Users/admin/...")` in mac
- Set the new working directory: `setwd("C:/Users/admin/...")` in windows
- Make sure the working directory changed: `getwd()`

B. Reading csv file:

```
file_need_to_read <- read.csv("C:/path/to/your/file.csv", header = TRUE)# in windows
file_need_to_read <- read.csv("/path/to/your/file.csv", header = TRUE)# in mac
```

```
cars_data <- read.csv("~/Desktop/STAT4101L_all_files/4101L-Fall-2023/cars.csv", header = TRUE)
head(cars_data, 2)
```

```
##      Make              Model Type Origin DriveTrain  MSRP Invoice
## 1 Subaru              Forester X Wagon      All 21445  19646
## 2 Toyota  Camry Solara SE V6 2dr Sedan   Asia   Front 21965  19819
##      EngineSize Cylinders Horsepower MPG_City MPG_Highway Weight Wheelbase Length
## 1           2.5         4         165      21         28   3090      99      175
## 2           3.3         6         225      20         29   3417     107     193
```

C. Reading xlsx file:

- Base R cannot read an excel sheet; however,
- it can be read using `read_excel` function from the “readxl” package.

Here we are going to read two excel datasets:

- **First install and load the library**

```
#install.packages("readxl") # install "readxl" package first if you have not installed yet
library(readxl) # load the library
```

```
# Syntax to read excel file
file_need_to_read <- read_excel("/path/to/your/file.xlsx")
```

Read Income data.

- This data can also be found in the Blackboard's dataset folder

```
library(readxl)
Income_data <- read_excel("~/Desktop/STAT4101L_all_files/4101L-Fall-2023/Course Materials/Section-3 Data")
print(Income_data)
```

```
## # A tibble: 10 x 4
##   storeID Sale_in_Thous Rent_in_Thous OtherIncome
##   <chr>      <dbl> <chr>      <chr>
## 1 12AR          165 7          NA
## 2 20AR          132 8           5
## 3 17AR          177 NA           4
## 4 11AR          128 NA          NA
## 5 26AR          137 5           3
## 6 18AR          199 NA           2
## 7 27AR          178 NA           6
## 8 25AR          104 6          NA
## 9 10AR          185 9           7
## 10 13AR          109 NA           1
```

D. Reading text file:

We use the `read.table()` function to import text data. It is important to determine how the data is separated and whether it includes a header.

- Tabular data

```
Class_marks <- read.table("/Users/suryalamichhane/Desktop/STAT4101L_all_files/Stat4101L-Rfiles/DataSets/Class_marks.txt")
head(Class_marks, 3) # print first 3 rows
```

```
##   Enrol.No. Maths Science English
## 1    A101    16     15     12
## 2    A102    16     17     11
## 3    A103    12     18     17
```

- Reading Text from a File (Uses: Processing text data)

```
Data_wrangling_text = readLines("/Users/suryalamichhane/Desktop/STAT4101L_all_files/Stat4101L-Rfiles/DataSets/Data_wrangling_text.txt")
print(Data_wrangling_text) #This text file was created by copying and pasting content from Google.
```

```
## [1] "Data wrangling, sometimes referred to as data munging, is the process of transforming and mapping"
```

- We can read text file into string using **readr** package

```
# Using readr package
library(readr)
Data_wrangling_text <- read_file("/Users/suryalamichhane/Desktop/STAT4101L_all_files/Stat4101L-Rfiles/d
Data_wrangling_text
```

```
## [1] "Data wrangling, sometimes referred to as data munging, is the process of transforming and mapping
```

```
# find number of characters
cat("Number of characters in our text documents: ", nchar(Data_wrangling_text), "\n")
```

```
## Number of characters in our text documents: 462
```

1.2 Previewing and Understanding Data

Previewing and understanding data are crucial steps in any data analysis process. In R, there are several functions and techniques we can use to get a sense of our data after the loading. Below are some common functions:

```
head(x, n = ..)
tail(x, n = ..)
str(x, n = ..)
View(x)
nrow(x)
ncol(x)
summary(x)
```

```
View(Income_data) #Invokes a spreadsheet style data viewer for the object
```

```
data(iris)
head(iris, n = 5) # top 5 observation of Iris flower data from R
```

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1          5.1          3.5          1.4          0.2 setosa
## 2          4.9          3.0          1.4          0.2 setosa
## 3          4.7          3.2          1.3          0.2 setosa
## 4          4.6          3.1          1.5          0.2 setosa
## 5          5.0          3.6          1.4          0.2 setosa
```

```
tail(iris, n = 5) # bottom 5 observation of Iris flower data from R
```

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 146          6.7          3.0          5.2          2.3 virginica
## 147          6.3          2.5          5.0          1.9 virginica
## 148          6.5          3.0          5.2          2.0 virginica
## 149          6.2          3.4          5.4          2.3 virginica
## 150          5.9          3.0          5.1          1.8 virginica
```

1.3. Missing Values

Handling missing values effectively is essential for maintaining the integrity of your dataset. Missing values should be handled appropriately. Here are the most common practices for dealing with missing values:

- Identify Missing Values: Before handling missing data, you must identify where and how many missing values exist.
- Remove Missing Values: Sometimes it is appropriate to remove rows or columns with missing data especially when data presents excessive missing values in a row or a column.
- Impute Missing Values: Replace missing values with plausible estimates.
 - Mean or Median Imputation (for numerical data)
 - Mode Imputation (for categorical data)
- Other practices are:
 - Interpolation: interpolate missing values based on surrounding data.
 - K-Nearest Neighbors (KNN) Imputation: Impute missing values based on the nearest neighbors in the dataset.
 - For categorical data, you can add a new category indicating missingness.
 - Use advanced statistical or machine learning models to predict and fill missing values.

1.4. Statistical Summary

```
# Summary statistics for a single column
summary(data$numeric_column)
# Compute specific metrics
mean(data$numeric_column, na.rm = TRUE)
median(data$numeric_column, na.rm = TRUE)
sd(data$numeric_column, na.rm = TRUE) # Standard deviation
var(data$numeric_column, na.rm = TRUE) # Variance
```

```
summary(iris)
```

```
##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
## Min.    :4.300   Min.    :2.000   Min.    :1.000   Min.    :0.100
## 1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
## Median :5.800   Median :3.000   Median :4.350   Median :1.300
## Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
## 3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
## Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
##      Species
## setosa    :50
## versicolor:50
## virginica :50
##
##
##
```

2. Subsetting Dataframes

- A subset can be created by using square brackets with specifying the row index (indices) and column index (indices)
- A dollar operator can be used to extract a column if we have column names
- Conditional filtering can be used to extract subset

2.1. Conditional filtering

a. Based on single condition

- Filter the dataset that corresponds to 'setosa' Species

```
setosa_data <- iris[iris$Species == 'setosa', ]  
head(setosa_data)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
## 1         5.1         3.5         1.4         0.2   setosa  
## 2         4.9         3.0         1.4         0.2   setosa  
## 3         4.7         3.2         1.3         0.2   setosa  
## 4         4.6         3.1         1.5         0.2   setosa  
## 5         5.0         3.6         1.4         0.2   setosa  
## 6         5.4         3.9         1.7         0.4   setosa
```

b. Based on multiple conditions

- Filter the dataset that corresponds to 'setosa' Species and Sepal.Length is between 5 and 6 cms.

```
filtered_setosa_data <- iris[iris$Species == 'setosa' &  
                             5 < iris$Sepal.Length & iris$Sepal.Length < 6, ]  
head(filtered_setosa_data)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
## 1         5.1         3.5         1.4         0.2   setosa  
## 6         5.4         3.9         1.7         0.4   setosa  
## 11        5.4         3.7         1.5         0.2   setosa  
## 15        5.8         4.0         1.2         0.2   setosa  
## 16        5.7         4.4         1.5         0.4   setosa  
## 17        5.4         3.9         1.3         0.4   setosa
```

2.2. The subset() Function

The subset() function can also be used to extract subsets of data if given conditions are met.

a. single condition

- Filter the dataset that corresponds to 'setosa' Species

```
setosa_data <- subset(iris, Species = 'setosa')  
head(setosa_data)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
## 1         5.1         3.5         1.4         0.2   setosa  
## 2         4.9         3.0         1.4         0.2   setosa  
## 3         4.7         3.2         1.3         0.2   setosa  
## 4         4.6         3.1         1.5         0.2   setosa  
## 5         5.0         3.6         1.4         0.2   setosa  
## 6         5.4         3.9         1.7         0.4   setosa
```

b. multiple conditions

- Filter the dataset that corresponds to 'setosa' Species and Sepal.Length is between 5 and 6 inches.

```
filtered_setosa_data <- subset(iris, Species == 'setosa' &
                               5 < Sepal.Length & Sepal.Length < 6)
head(filtered_setosa_data)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1           5.1         3.5         1.4         0.2   setosa
## 6           5.4         3.9         1.7         0.4   setosa
## 11          5.4         3.7         1.5         0.2   setosa
## 15          5.8         4.0         1.2         0.2   setosa
## 16          5.7         4.4         1.5         0.4   setosa
## 17          5.4         3.9         1.3         0.4   setosa
```

2.3. Select, drop or add variables in Dataframes

a. Select variables using column index

- Select Sepal.Width and Petal.Width from iris data

```
# select using the columns indices
colnames <- c("Sepal.Length", "Sepal.Width", "Petal.Length", "Petal.Width", "Species")
iris_new1 <- iris[, c(2, 4)]
head(iris_new1)
```

```
##      Sepal.Width Petal.Width
## 1           3.5         0.2
## 2           3.0         0.2
## 3           3.2         0.2
## 4           3.1         0.2
## 5           3.6         0.2
## 6           3.9         0.4
```

b. Select variables using names

- Select Sepal.Width and Petal.Width from iris data

```
# select using the variable names
iris_new2 <- iris[, c("Sepal.Width", "Petal.Width")]
head(iris_new2)
```

```
##      Sepal.Width Petal.Width
## 1           3.5         0.2
## 2           3.0         0.2
## 3           3.2         0.2
## 4           3.1         0.2
## 5           3.6         0.2
## 6           3.9         0.4
```

c. Select variables using subset

- Select Sepal.Width and Petal.Width from iris data

```
# select using the subset
iris_new3 <- subset(iris, select = c("Sepal.Width", "Petal.Width"))
head(iris_new3)
```

```
##   Sepal.Width Petal.Width
## 1         3.5         0.2
## 2         3.0         0.2
## 3         3.2         0.2
## 4         3.1         0.2
## 5         3.6         0.2
## 6         3.9         0.4
```

d. Dropping variables using negative column index

- Drop Sepal.Width and Petal.Width from iris data

```
# can be dropped using -ve sign of the columns index
iris_new4 <- iris[, - c(2, 4)]
head(iris_new4)
```

```
##   Sepal.Length Petal.Length Species
## 1         5.1         1.4  setosa
## 2         4.9         1.4  setosa
## 3         4.7         1.3  setosa
## 4         4.6         1.5  setosa
## 5         5.0         1.4  setosa
## 6         5.4         1.7  setosa
```

e. Dropping variables using subset

- Drop Sepal.Width and Petal.Width from iris data

```
iris_new5 <- subset(iris, select = - c(Sepal.Width, Petal.Width))
head(iris_new5, 2)
```

```
##   Sepal.Length Petal.Length Species
## 1         5.1         1.4  setosa
## 2         4.9         1.4  setosa
```

- f. **Add new variables in data frame** <div class="alert alert-block alert-info", style="margin-top: 20px"> Create a categorical variable 'Sepal_Len_cat' based on following rule and add it to the iris data. :

```
Sepal.Length <= 5, category="low"
5 < Sepal.Length <= 6, category="low_mid"
6 < Sepal.Length <= 7, category="high_mid"
7 < Sepal.Length <= 8, category="high",
```


- Let's create a variable using loop and conditions
- We use `cut()` function

```
var1 <- iris$Sepal.Length
cut_off <- c(0, 5, 6, 7, 8)
category <- c("low", "low_mid", "high_mid", "high")
Sepal_Len_cat <- cut(var1, breaks = cut_off, labels = category)
iris_new <- cbind(iris, Sepal_Len_cat)
head(iris_new)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species Sepal_Len_cat
## 1         5.1         3.5         1.4         0.2  setosa      low_mid
## 2         4.9         3.0         1.4         0.2  setosa        low
## 3         4.7         3.2         1.3         0.2  setosa        low
## 4         4.6         3.1         1.5         0.2  setosa        low
## 5         5.0         3.6         1.4         0.2  setosa        low
## 6         5.4         3.9         1.7         0.4  setosa      low_mid
```

3. Sorting data in R

a. Sorting based on single column

- use `order()` function
- `order()` function returns the indices of the entries in desired order
- Syntax : `order(x, decreasing = FALSE)`

```
sx <- c(3, 4, 2, 4)
order(sx)
```

```
## [1] 3 1 2 4
```

In the above example shows smallest entry is 2 which is in 3rd position, 2nd smallest entry is 3 which is in first position and so on. We now order the data based on their positions that we found above.

```
sx <- c(3, 4, 2, 4)
order(sx, decreasing = TRUE)
```

```
## [1] 2 4 1 3
```

More Practice of order function

- rearrange the `iris_new` data in ascending order of `Sepal.Length` as in above

```
sorted_iris <- iris_new[order(iris_new$Sepal.Length, decreasing = FALSE), ]
head(sorted_iris)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width Species Sepal_Len_cat
## 14          4.3         3.0         1.1         0.1  setosa          low
## 9           4.4         2.9         1.4         0.2  setosa          low
## 39          4.4         3.0         1.3         0.2  setosa          low
## 43          4.4         3.2         1.3         0.2  setosa          low
## 42          4.5         2.3         1.3         0.3  setosa          low
## 4           4.6         3.1         1.5         0.2  setosa          low
```

- rearrange the iris_new data in descending order of Sepal.Length

```
sorted_iris <- iris_new[order(iris_new$Sepal.Length, decreasing = TRUE), ]
head(sorted_iris)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width  Species Sepal_Len_cat
## 132          7.9         3.8         6.4         2.0 virginica          high
## 118          7.7         3.8         6.7         2.2 virginica          high
## 119          7.7         2.6         6.9         2.3 virginica          high
## 123          7.7         2.8         6.7         2.0 virginica          high
## 136          7.7         3.0         6.1         2.3 virginica          high
## 106          7.6         3.0         6.6         2.1 virginica          high
```

b. Sorting based on multiple column

- rearrange the iris_new data in ascending order of Sepal.Length, Sepal.Width and Species type

```
sorted_iris <- iris_new[order(iris_new$Sepal.Length, iris_new$Sepal.Width, iris_new$Species,
                             decreasing = FALSE), ]
head(sorted_iris)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width Species Sepal_Len_cat
## 14          4.3         3.0         1.1         0.1  setosa          low
## 9           4.4         2.9         1.4         0.2  setosa          low
## 39          4.4         3.0         1.3         0.2  setosa          low
## 43          4.4         3.2         1.3         0.2  setosa          low
## 42          4.5         2.3         1.3         0.3  setosa          low
## 4           4.6         3.1         1.5         0.2  setosa          low
```

c. Decoding a variable

Let's decode the variable "Sepal_Len_cat" as "low" = 1, "low_mid" = 2, "high_mid" = 3, "high" = 4.

```
# make sure levels are in order
iris_new$Sepal_Len_cat <- factor(iris_new$Sepal_Len_cat, levels = c("low", "low_mid", "high_mid", "high"))

# Decoding the levels as 1, 2, 3, 4
iris_new$Sepal_Len_decoded <- as.numeric(iris_new$Sepal_Len_cat)

# View the result
head(iris_new)
```

```
## Sepal.Length Sepal.Width Petal.Length Petal.Width Species Sepal_Len_cat
## 1          5.1          3.5          1.4          0.2 setosa      low_mid
## 2          4.9          3.0          1.4          0.2 setosa        low
## 3          4.7          3.2          1.3          0.2 setosa        low
## 4          4.6          3.1          1.5          0.2 setosa        low
## 5          5.0          3.6          1.4          0.2 setosa        low
## 6          5.4          3.9          1.7          0.4 setosa      low_mid
## Sepal_Len_decoded
## 1          2
## 2          1
## 3          1
## 4          1
## 5          1
## 6          2
```

4. Merging of Data Tables

Data merging is essential in data analysis, data science, and database management because it allows combining information from multiple sources to create a comprehensive dataset.

- Data is usually merged based on the primary key,
- Primary key is the set of data columns that are common in dataframes.
- “merge()” function in R combines two data tables at a time.

Syntax()

merge(x,y, ...), where x, y are data frames to be coerced to one.

Examples (Merging datasets)

- Table-1: Income data of 10 variuos store
- Table-2: Expense data of 10 variuos store

Now we want to find the details of from these two dataset.

```
library(readxl)
Income_data <- read_excel("~/Desktop/STAT4101L_all_files/4101L-Fall-2023/Course Materials/Section-3 Data")
print(Income_data)
```

```
## # A tibble: 10 x 4
##   storeID Sale_in_Thous Rent_in_Thous OtherIncome
##   <chr>      <dbl> <chr>      <chr>
## 1 12AR          165 7          NA
## 2 20AR          132 8           5
## 3 17AR          177 NA           4
## 4 11AR          128 NA          NA
## 5 26AR          137 5           3
## 6 18AR          199 NA           2
## 7 27AR          178 NA           6
## 8 25AR          104 6          NA
## 9 10AR          185 9           7
## 10 13AR          109 NA           1
```

```
Expense_data <- read_excel("~/Desktop/STAT4101L_all_files/4101L-Fall-2023/Course Materials/Section-3 Data/Expense_data.xlsx")
print(Expense_data)
```

```
## # A tibble: 10 x 4
##   storeID Purchase_in_Thous EmployeeCost OtherCost
##   <chr>          <dbl>          <dbl>      <dbl>
## 1 19AR             120             10         15
## 2 29AR             107             10         11
## 3 27AR              98             11         12
## 4 18AR              86             15         12
## 5 13AR              81             14         14
## 6 20AR             103             15         15
## 7 30AR             138             11         11
## 8 14AR             128             14         11
## 9 25AR             127             14         12
## 10 26AR            135             15         10
```

```
merge(Income_data, Expense_data, by = 'storeID')
```

```
##   storeID Sale_in_Thous Rent_in_Thous OtherIncome Purchase_in_Thous
## 1   13AR           109           NA             1             81
## 2   18AR           199           NA             2             86
## 3   20AR           132            8             5            103
## 4   25AR           104            6            NA            127
## 5   26AR           137            5             3            135
## 6   27AR           178           NA             6             98
##   EmployeeCost OtherCost
## 1            14         14
## 2            15         12
## 3            15         15
## 4            14         12
## 5            15         10
## 6            11         12
```

In the above two tables: Both tables have same common id, when the id names are different we use by.x, by.y

```
class_info <- data.frame(ID = c(001, 002, 005),
                          Name = c("Alex", "Mia", "Sam"))

test_score <- data.frame(stID = c(001, 001, 002, 003),
                          Course = c("Math", "Hist", "Math", "Math"),
                          Score = c(73, 82, 88, 80))

merge(class_info, test_score, by.x = 'ID', by.y = 'stID')
```

```
##   ID Name Course Score
## 1  1 Alex  Math    73
## 2  1 Alex  Hist    82
## 3  2 Mia   Math    88
```

Types of Merge

1. Inner merge: An inner merge combines two dataframes to include only the rows with matching primary keys in both dataframes. The resulting dataframe will have the following characteristics:

- Only matching rows are included: Rows are included in the resulting dataframe only if the primary key exists in both dataframes.
- No unmatched rows: Any row from either dataframe that does not have a match in the other dataframe is excluded from the result.
- Columns from both dataframes are combined for matching rows: For rows where the primary key matches, the columns from both dataframes are merged into a single row.

This ensures that the resulting dataframe contains only data that is common to both dataframes.

```
merge(class_info, test_score, by.x = 'ID', by.y = 'stID', all = FALSE)
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
```

2. Left merge: A left merge combines two dataframes to include all rows from the first dataframe (“x”) and only the matching rows from the second dataframe (“y”). The resulting dataframe will have the following characteristics:

- a. All rows with common primary keys are included: Rows that have matching primary key values in both dataframes are combined into the resulting dataframe.
- b. All rows from the first dataframe are included, even if there is no match in the second dataframe: If a primary key exists in the first dataframe but not in the second, that row from the first dataframe will still appear in the result.
- c. Missing values (NA) will appear for unmatched columns from the second dataframe: For rows where a primary key is present only in the first dataframe and not in the second, the columns from the second dataframe will be filled with NA values.

```
merge(class_info, test_score, by.x = 'ID', by.y = 'stID', all.x = TRUE)
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
## 4  5 Sam    <NA>    NA
```

3. Right merge: A right merge combines two dataframes to include all rows from the second dataframe (“y”) and only the matching rows from the first dataframe (“x”).

```
merge(class_info, test_score, by.x = 'ID', by.y = 'stID', all.y = TRUE)
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
## 4  3 <NA>   Math    80
```

4. Outer merge: An outer merge combines two dataframes to include all rows from both dataframes (“x” and “y”), regardless of whether there is a match in their primary keys. The resulting dataframe will have the following characteristics:

- a. All rows from both dataframes are included: Every row from both dataframes is part of the resulting dataframe, even if there is no match in the primary keys.
- b. Matching rows are merged: If a primary key exists in both dataframes, the rows are combined, and data from both dataframes is included.
- c. Missing values (NA) appear for unmatched rows:
 - For rows present only in the first dataframe (x) and not in the second (y), the columns from the second dataframe will be filled with NA.
 - Similarly, for rows present only in the second dataframe (y) and not in the first (x), the columns from the first dataframe will be filled with NA.

This ensures that no data is lost from either dataframe, and all unique primary keys are represented in the resulting dataframe.

```
Merged_data <- merge(class_info, test_score, by.x = 'ID', by.y = 'stID', all = TRUE)
Merged_data
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
## 4  3 <NA>   Math    80
## 5  5 Sam    <NA>    NA
```

5. dplyr Package for Data Manipulation

- Install package: `install.packages('dplyr')`
- Most popular function in *dplyr* Package

select: select variables and find subset filter: find subset based on conditional filtering mutate: create a new variables or modify existing variables arrange: sorting and ordering groupby: aggregating variables summarize: aggregating variables and finding aggregated values

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

5.1. select

```
select(Merged_data, ID, Name, Score)
```

```
##   ID Name Score
## 1  1 Alex    73
## 2  1 Alex    82
## 3  2 Mia    88
## 4  3 <NA>    80
## 5  5 Sam     NA
```

- Using Pyping operator: %>%

```
Merged_data %>%
  select(ID, Name, Score) %>%
  head(n = 3) # print top 3
```

```
##   ID Name Score
## 1  1 Alex    73
## 2  1 Alex    82
## 3  2 Mia    88
```

```
# Drop the columns of the dataframe and print top 3
select(Merged_data, -c("Name", "Score")) %>% head(n = 3)
```

```
##   ID Course
## 1  1  Math
## 2  1  Hist
## 3  2  Math
```

5.2. filter

```
filter(Merged_data, Course == "Math")
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  2 Mia   Math    88
## 3  3 <NA>  Math    80
```

- Filter the Courses Math and Hist and student is Alex

```
filter(Merged_data, Name == "Alex", Course %in% c("Math", "Hist")) %>% head()
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
```

```
Merged_data %>% filter(Name == "Alex", Course %in% c("Math", "Hist")) %>% head()
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
```

5.3. mutate (Uses: Create a new variable)

Example: Create a new variable individual_average

```
cars_data <- read.csv("~/Desktop/STAT4101L_all_files/4101L-Fall-2023/cars.csv", header = TRUE)
head(cars_data)
```

```
##      Make              Model Type Origin DriveTrain  MSRP Invoice
## 1  Subaru      Forester X Wagon      All 21445  19646
## 2 Toyota  Camry Solara SE V6 2dr Sedan   Asia   Front 21965  19819
## 3 Suzuki      Aerio LX 4dr Sedan   Asia   Front 14500  14317
## 4  Dodge      Dakota Club Cab Truck    USA    Rear 20300  18670
## 5  Mazda      Mazda3 s 4dr Sedan   Asia   Front    NA  15922
## 6 Infiniti    G35 Sport Coupe 2dr Sedan   Asia    Rear 29795  27536
##   EngineSize Cylinders Horsepower MPG_City MPG_Highway Weight Wheelbase Length
## 1         2.5         4        165      21         28   3090      99      175
## 2         3.3         6        225      20         29   3417     107     193
## 3         2.3         4        155      25         31   2676      98     171
## 4         3.7         6        210      16         22   3829     131     219
## 5         2.3         4        160      25         31   2762     104     179
## 6         3.5         6        280      18         26   3416     112     182
```

c. mutate

- Create a new variable $\text{avg_mpg} = (\text{MPG_City} + \text{MPG_Highway})/2$

```
new_cars_data <- cars_data %>%
  mutate(avg_mpg = (MPG_City + MPG_Highway)/2)

new_cars_data %>% head(3)
```

```
##      Make              Model Type Origin DriveTrain  MSRP Invoice
## 1 Subaru      Forester X Wagon      All 21445  19646
## 2 Toyota  Camry Solara SE V6 2dr Sedan   Asia   Front 21965  19819
## 3 Suzuki      Aerio LX 4dr Sedan   Asia   Front 14500  14317
##   EngineSize Cylinders Horsepower MPG_City MPG_Highway Weight Wheelbase Length
## 1         2.5         4        165      21         28   3090      99      175
## 2         3.3         6        225      20         29   3417     107     193
```



```
## 3      2.3      4      155      25      31  2676      98  171
##   avg_mpg
## 1    24.5
## 2    24.5
## 3    28.0
```

5.4. arrange (Uses: to arrange the dataset in some specific order)

- Sort the dataset in increasing order of Make and decreasing order of MSRP.

```
cars_data_new2 = cars_data %>%
  arrange(Make, desc(MSRP))

head(cars_data_new2)
```

```
##   Make      Model   Type Origin DriveTrain  MSRP Invoice
## 1 Acura  NSX coupe 2dr manual S Sports   Asia      Rear 89765  79978
## 2 Acura  3.5 RL w/Navigation 4dr Sedan   Asia      Front 46100  41100
## 3 Acura      3.5 RL 4dr Sedan   Asia      Front 43755  39014
## 4 Acura      TL 4dr Sedan   Asia      Front 33195  30299
## 5 Acura      RSX Type S 2dr Sedan   Asia      Front 23820  21761
## 6 Acura      TSX 4dr Sedan   Asia      Front    NA  24647
##   EngineSize Cylinders Horsepower MPG_City MPG_Highway Weight Wheelbase Length
## 1      3.2      6      290      17      24   3153      100   174
## 2      3.5      6      225      18      24   3893      115   197
## 3      3.5      6      225      18      24   3880      115   197
## 4      3.2      6      270      20      28   3575      108   186
## 5      2.0      4      200      24      31   2778      101   172
## 6      2.4      4      200      22      29   3230      105   183
```

5.5. Data Aggregation

- Aggregation includes process of summarizing group wise the data based on levels of a factor column.
- Aggregation in R can be done using functions *aggregate()* and *tapply()*.
- summarise() and group_by()

```
summarise(group_by(cars_data, Type),
  mean_price = mean(MSRP, na.rm = TRUE), count = n() )
```

```
## # A tibble: 6 x 3
##   Type    mean_price count
##   <chr>      <dbl> <int>
## 1 Hybrid    20325      3
## 2 SUV      34447.     60
## 3 Sedan    29716.    262
## 4 Sports   53793.     49
## 5 Truck    22967.     24
## 6 Wagon    29188.     30
```

5.6 Decoding data using *dplyr* package

```
# Remove the 'Sepal_Len_cat1' column and recode 'Sepal_Len_cat'
iris_new <- iris_new %>%
  mutate(Sepal_Len_cat = case_when(
    Sepal_Len_cat == "low" ~ 1,
    Sepal_Len_cat == "low_mid" ~ 2,
    Sepal_Len_cat == "high_mid" ~ 3,
    Sepal_Len_cat == "high" ~ 4
  ))

# Display the first few rows of the updated dataset
head(iris_new)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species Sepal_Len_cat
## 1         5.1         3.5         1.4         0.2   setosa             2
## 2         4.9         3.0         1.4         0.2   setosa             1
## 3         4.7         3.2         1.3         0.2   setosa             1
## 4         4.6         3.1         1.5         0.2   setosa             1
## 5         5.0         3.6         1.4         0.2   setosa             1
## 6         5.4         3.9         1.7         0.4   setosa             2
##   Sepal_Len_decoded
## 1                 2
## 2                 1
## 3                 1
## 4                 1
## 5                 1
## 6                 2
```

5.7 Merging data using *dplyr*

```
head(class_info)
```

```
##   ID Name
## 1  1 Alex
## 2  2 Mia
## 3  5 Sam
```

```
test_score
```

```
##   stID Course Score
## 1    1   Math    73
## 2    1   Hist    82
## 3    2   Math    88
## 4    3   Math    80
```

```
Left_join_data <- left_join(class_info, test_score, by = c("ID"="stID"))
Left_join_data
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
## 4  5 Sam    <NA>    NA
```

```
Inner_join_data <- inner_join(class_info, test_score, by = c("ID"="stID"))
Inner_join_data
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
```

```
full_join_data <- full_join(class_info, test_score, by = c("ID"="stID"))
full_join_data
```

```
##   ID Name Course Score
## 1  1 Alex   Math    73
## 2  1 Alex   Hist    82
## 3  2 Mia    Math    88
## 4  5 Sam    <NA>    NA
## 5  3 <NA>   Math    80
```

6. Data Reshaping using tidyr

The tidyr package in R is powerful for data wrangling and analysis. It is commonly used in conjunction with dplyr for tidying, transforming, and analyzing data before applying statistical models or visualizations. Below are key data analysis workflows using tidyr.

6.1. Load Required Library

```
library(tidyr)
```

Example Data: Student Scores Dataset

```
# Sample data
df <- data.frame(
  ID = c(1, 1, 2, 2, 3, 3, 4),
  Name = c("Alex", "Alex", "Mia", "Mia", "Sam", "Sam", NA),
  Course = c("Math", "Hist", "Math", "Hist", "Math", "Hist", "Math"),
  Score = c(90, 85, 88, 92, 78, NA, 95)
)

print(df)
```

```
##   ID Name Course Score
## 1  1 Alex   Math    90
## 2  1 Alex   Hist    85
```

```
## 3  2  Mia   Math   88
## 4  2  Mia   Hist   92
## 5  3  Sam   Math   78
## 6  3  Sam   Hist   NA
## 7  4 <NA>   Math   95
```

6.2. Data Cleaning and Tidying with tidyr

Handling Missing Values: `replace_na()` & `drop_na()` We have missing values in the Name and Score columns. We can either remove them or replace them.

```
df_clean <- df %>%
  replace_na(list(Name = "Unknown", Score = 0)) # Replace missing names with "Unknown" and scores with 0
print(df_clean)
```

```
##   ID   Name Course Score
## 1  1   Alex   Math    90
## 2  1   Alex   Hist    85
## 3  2    Mia   Math    88
## 4  2    Mia   Hist    92
## 5  3    Sam   Math    78
## 6  3    Sam   Hist     0
## 7  4 Unknown   Math    95
```

```
df_no_na <- df %>%
  drop_na() # remove rows with missing value
print(df_no_na)
```

```
##   ID Name Course Score
## 1  1 Alex   Math    90
## 2  1 Alex   Hist    85
## 3  2 Mia   Math    88
## 4  2 Mia   Hist    92
## 5  3 Sam   Math    78
```

6.3. Reshaping Data: `pivot_wider()` & `pivot_longer()`

After tidying, we may need to reshape the data.

```
df_wide <- df_clean %>%
  pivot_wider(names_from = Course, values_from = Score) # Convert Long to Wide Format
print(df_wide)
```

```
## # A tibble: 4 x 4
##       ID Name    Math Hist
##   <dbl> <chr>  <dbl> <dbl>
## 1     1 Alex     90     85
## 2     2 Mia      88     92
## 3     3 Sam      78      0
## 4     4 Unknown  95     NA
```

```
df_long <- df_wide %>%
  pivot_longer(cols = c(Math, Hist), names_to = "Course", values_to = "Score") #Convert Wide to Long Fo
print(df_long)
```

```
## # A tibble: 8 x 4
##       ID Name    Course Score
##   <dbl> <chr>   <chr>  <dbl>
## 1     1 Alex    Math    90
## 2     1 Alex    Hist    85
## 3     2 Mia     Math    88
## 4     2 Mia     Hist    92
## 5     3 Sam     Math    78
## 6     3 Sam     Hist     0
## 7     4 Unknown Math    95
## 8     4 Unknown Hist    NA
```